

Serverless Machine Learning on Modern Hardware

Patrick Stuedi
IBM Research

#Res6SAIS

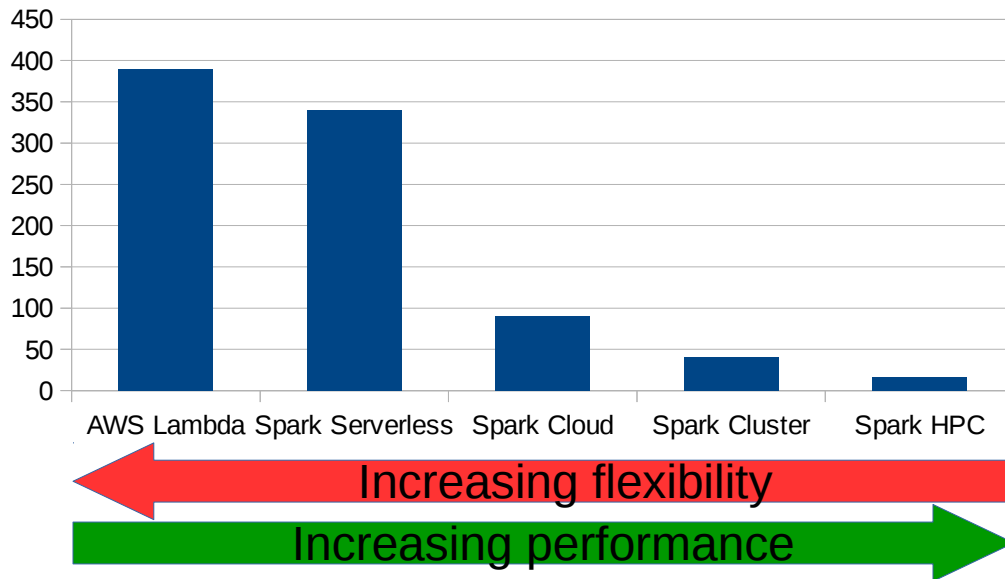
Serverless Computing



- No need to setup/manage a cluster
- Automatic, dynamic and fine-grained scaling
- Sub-second billing
- AWS Lambda, Google Cloud Functions, Azure Functions, Databricks Serverless

Challenge: Performance


Example:
Sorting 100GB



Challenge: Performance


Example:
Sorting 100GB

64 worker
cores




Challenge: Performance

Example:
Sorting 100GB




Challenge: Performance

Example:
Sorting 100GB



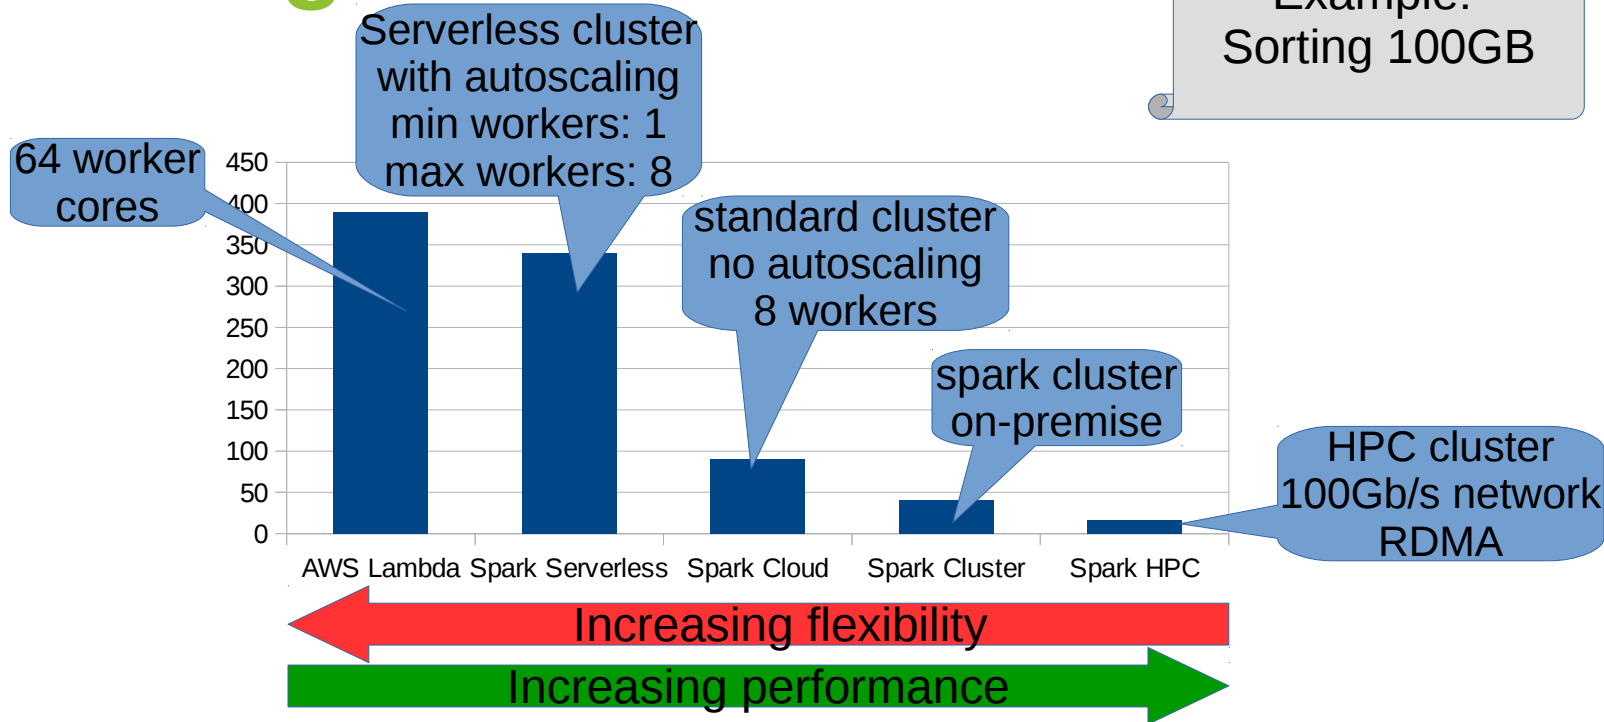
Challenge: Performance

Example:
Sorting 100GB



Challenge: Performance

Example:
Sorting 100GB



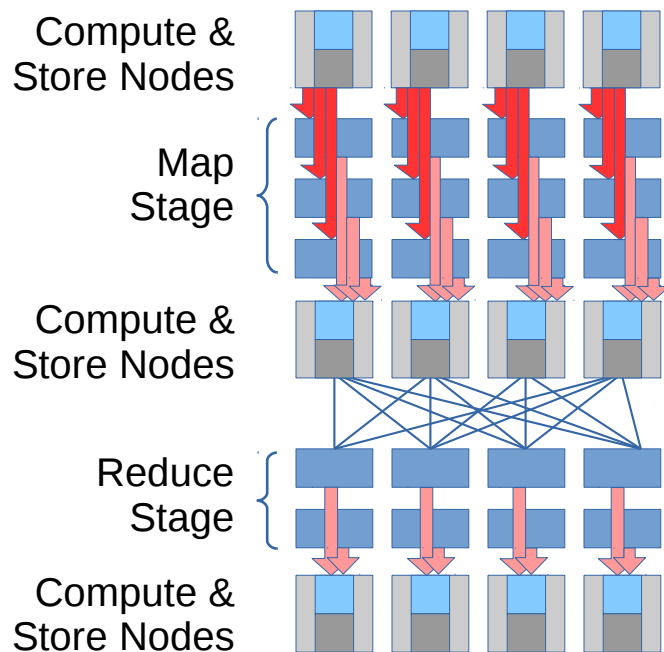
Why is it so hard?

- **Scheduler:** when to best add/remove resources?
- **Container startup:** may have to dynamically spin up containers per function
- **Storage:** input data needs to be fetched from remote storage (e.g., S3)
 - As opposed to compute-local storage, e.g., HDFS
- **Data sharing:** intermediate needs to be temporarily stored on remote storage (S3, Redis)
 - Affects operations like shuffle, broadcast, etc.,

Why is it so hard?


- **Scheduler:** when to best add/remove resources?
- **Container startup:** may have to dynamically spin up containers per function
- **Storage:** input data needs to be fetched from remote storage (e.g., S3)
 - As opposed to compute-local storage, e.g., HDFS
- **Data sharing:** intermediate needs to be temporarily stored on remote storage (S3, Redis)
 - Affects operations like shuffle, broadcast, etc.,

Example: MapReduce (Cluster)




data is mostly
written and
read locally

Example: MapReduce (Serverless)




I/O Overhead: Sorting 100GB



Input/output and shuffle overheads are significantly higher when data is stored remotely

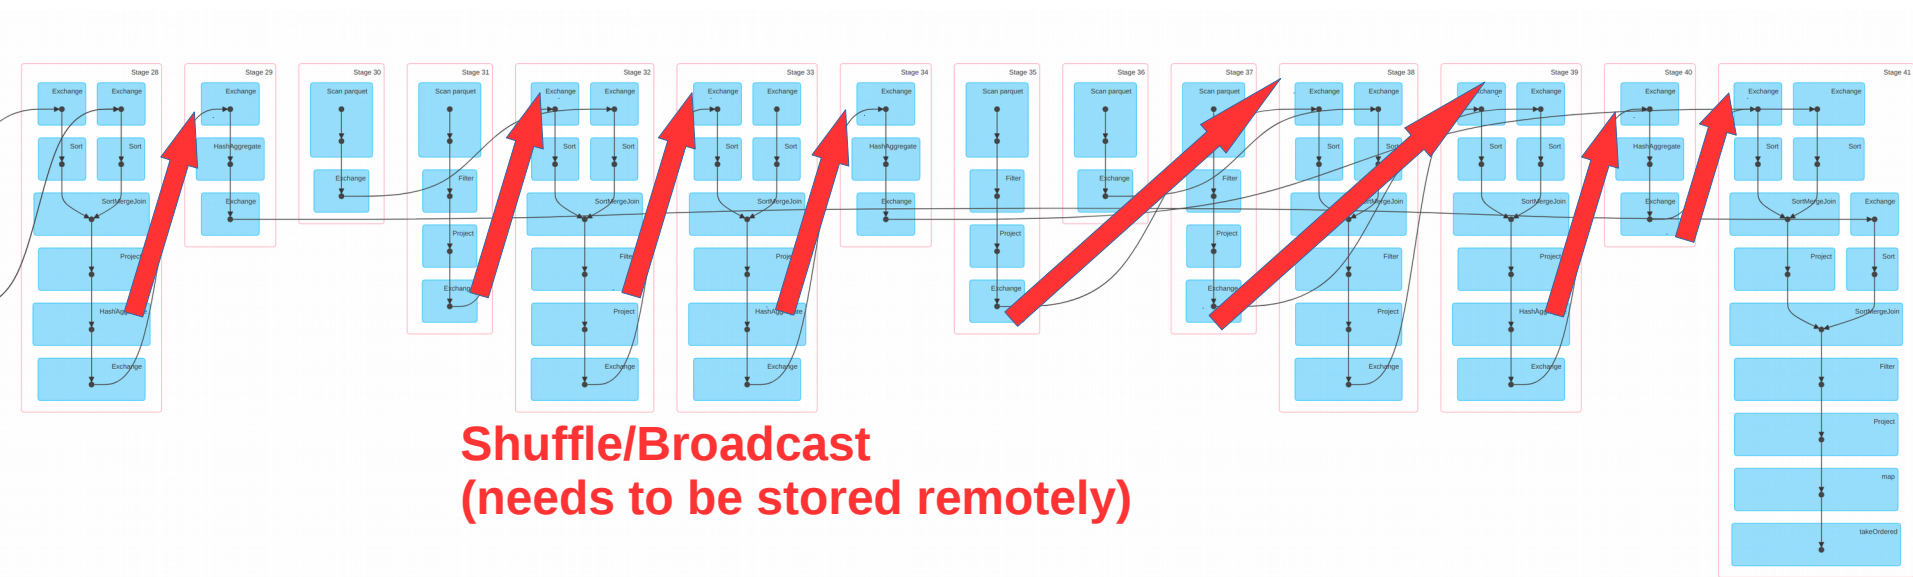
What about other workloads?

Example: SQL, Query 77 / TPC-DS benchmark



What about other workloads?

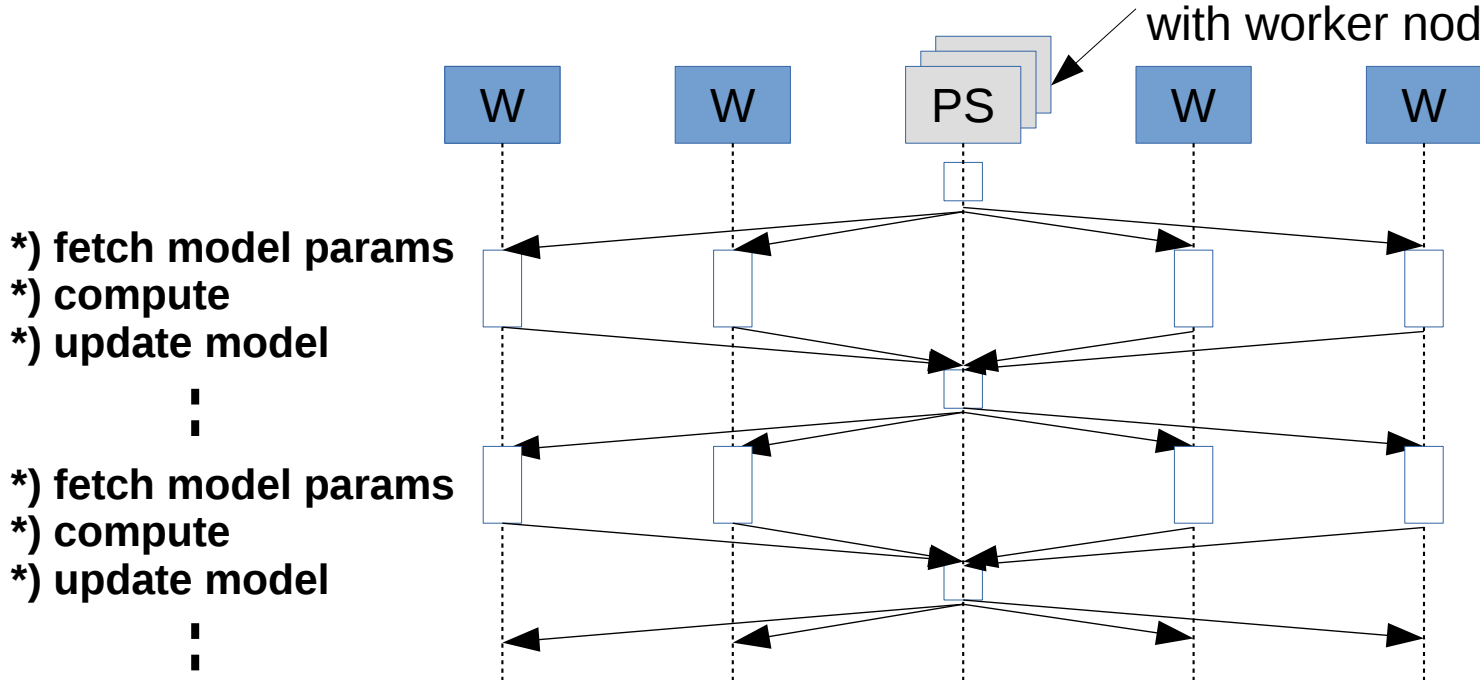
Example: SQL, Query 77 / TPC-DS benchmark



What about other workloads?

Example: Iterative ML (e.g., linear regression)

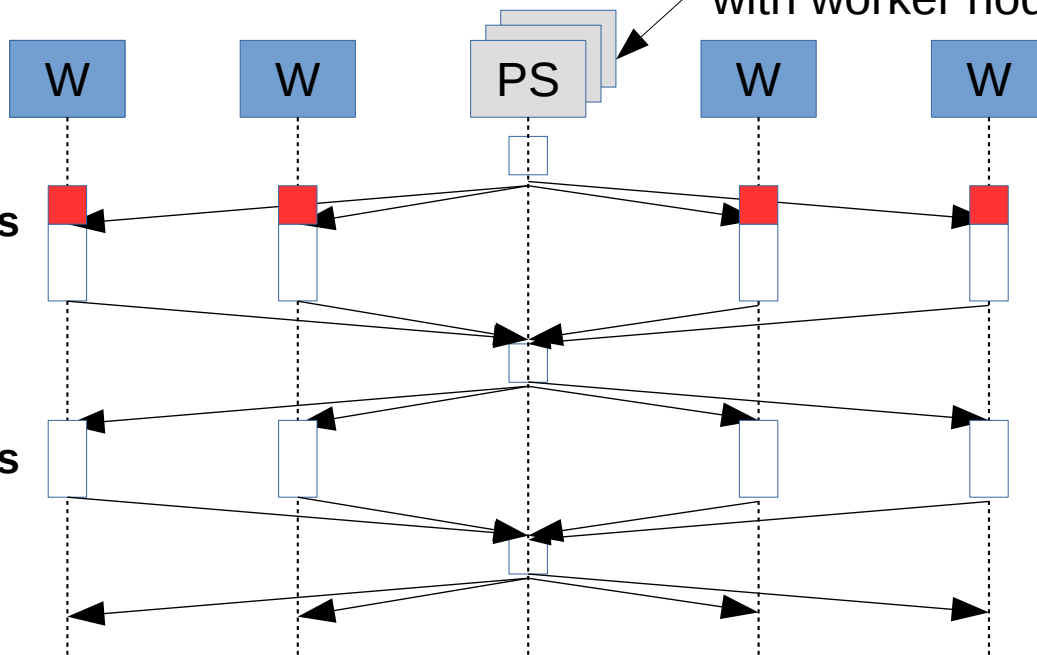
could be co-located
with worker nodes



What about other workloads?

Example: Iterative ML (e.g., linear regression)

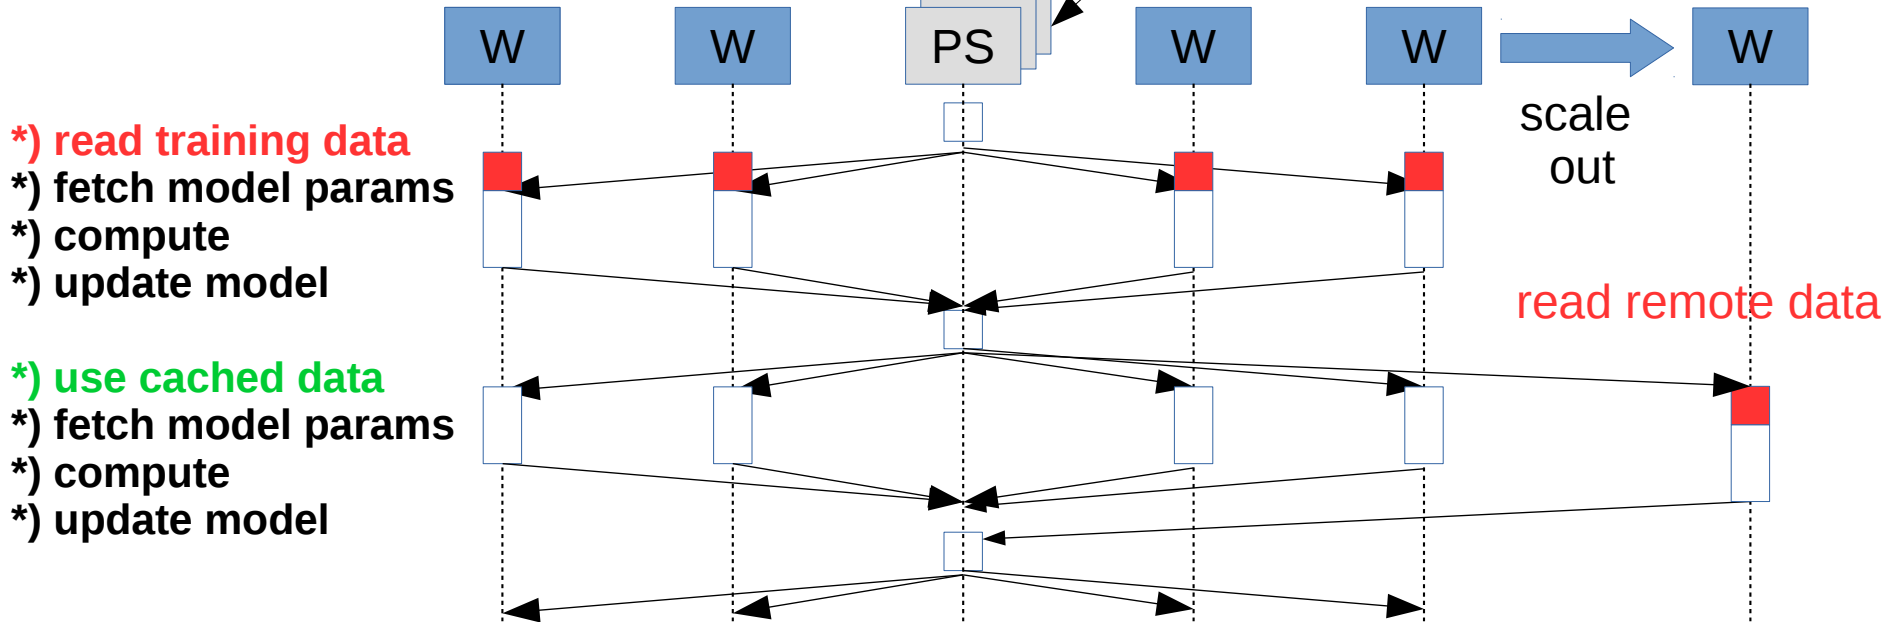
could be co-located
with worker nodes



What about other workloads?

Example: Iterative ML (e.g., linear regression)

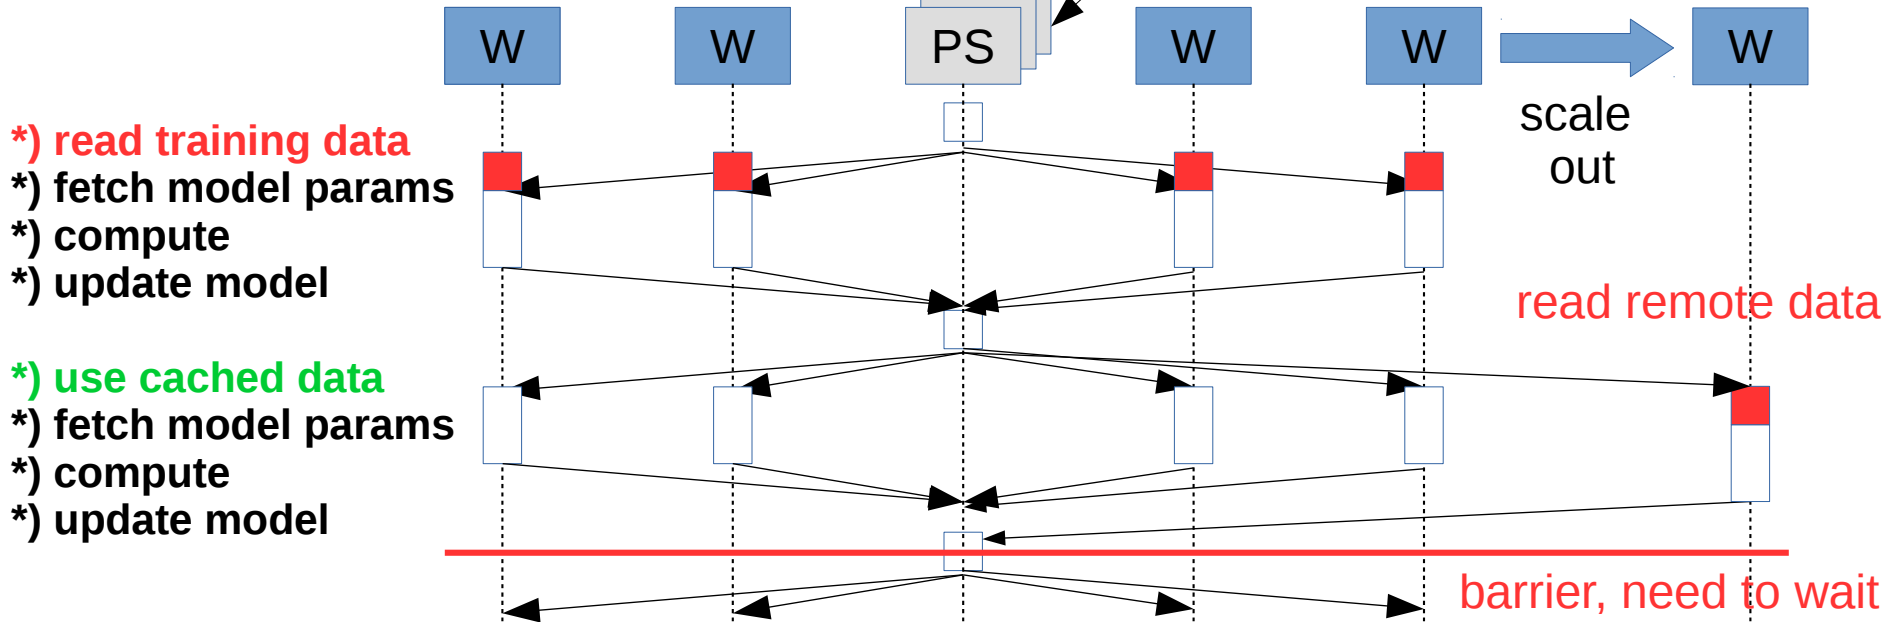
~~could be co-located
with worker nodes~~ Needs to be
remote



What about other workloads?

Example: Iterative ML (e.g., linear regression)

~~could be co-located with worker nodes~~ Needs to be remote



Can we..

- Use Spark to run such workloads in a serverless fashion?
 - Dynamic scaling of compute nodes as jobs are running
 - No cluster configuration
 - No startup time
- Reduce the performance overheads to a minimum?

Design Options

- **Scheduling:**

- 1 Use serverless framework to schedule executors
- 2 Use serverless framework to schedule tasks
- 3 Enable Spark to dynamically scale up and down executors

- **Intermediate data:**

- 1 Executors cooperate with scheduler to flush data remotely
- 2 Consequently store all intermediate state remotely

Design Options

- **Scheduling:**

High startup
Latency!

- 1 Use serverless framework to schedule executors
- 2 Use serverless framework to schedule tasks
- 3 Enable Spark to dynamically scale up and down executors

- **Intermediate data:**

- 1 Executors cooperate with scheduler to flush data remotely
- 2 Consequently store all intermediate state remotely

Design Options

- **Scheduling:**

- ① Use serverless framework to schedule executors
- ② Use serverless framework to schedule tasks
- ③ Enable Spark to dynamically scale up and down executors

High startup
Latency!

Slow!

- **Intermediate data:**

- ① Executors cooperate with scheduler to flush data remotely
- ② Consequently store all intermediate state remotely

Design Options

- **Scheduling:**

- 1 Use serverless framework to schedule executors
- 2 Use serverless framework to schedule tasks
- 3 Enable Spark to dynamically scale up and down executors

High startup
Latency!

Slow!

- **Intermediate data:**

- 1 Executors cooperate with scheduler to flush data remotely
- 2 Consequently store all intermediate state remotely

Design Options

- **Scheduling:**

- 1 Use serverless framework to schedule executors
- 2 Use serverless framework to schedule tasks
- 3 Enable Spark to dynamically scale up and down executors

High startup
Latency!

Slow!

- **Intermediate data:**

- 1 Executors cooperate with scheduler to flush data remotely
- 2 Consequently store all intermediate state remotely

Complex!

Design Options

- **Scheduling:**

- 1 Use serverless framework to schedule executors
- 2 Use serverless framework to schedule tasks
- 3 Enable Spark to dynamically scale up and down executors

High startup
Latency!

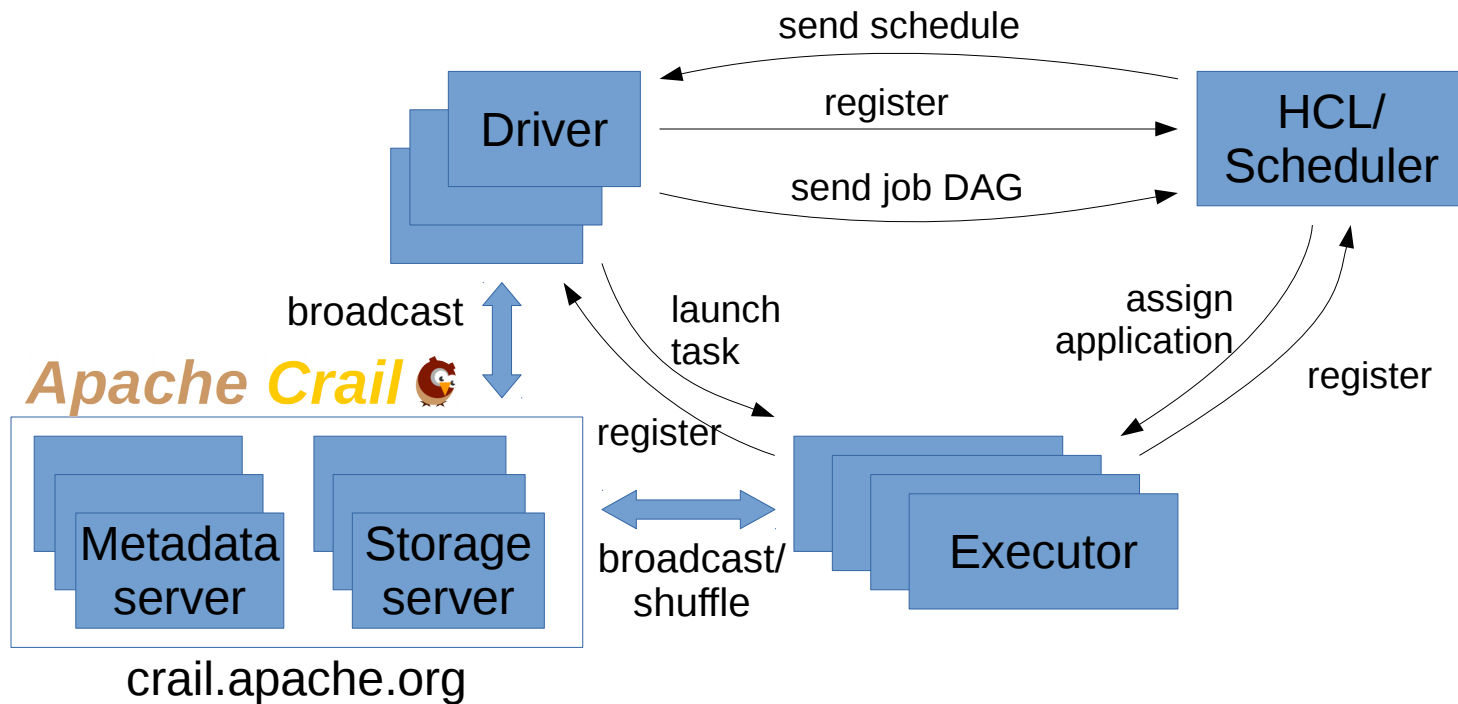
Slow!

- **Intermediate data:**


- 1 Executors cooperate with scheduler to flush data remotely
- 2 Consequently store all intermediate state remotely

Complex!


Architecture Overview




Architecture Overview



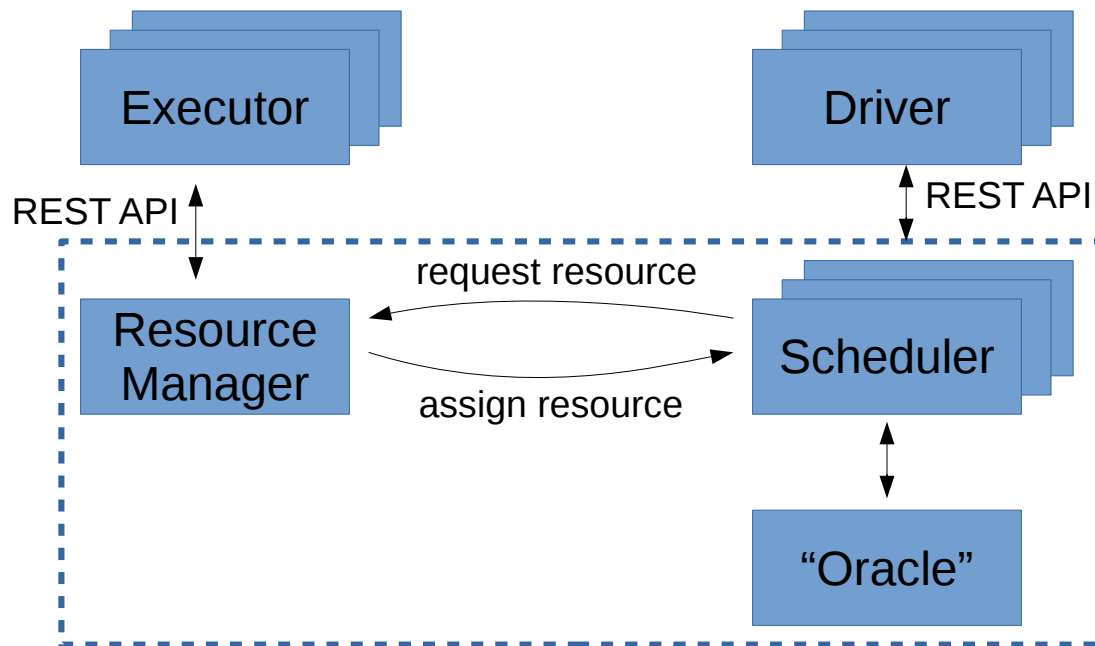
Architecture Overview



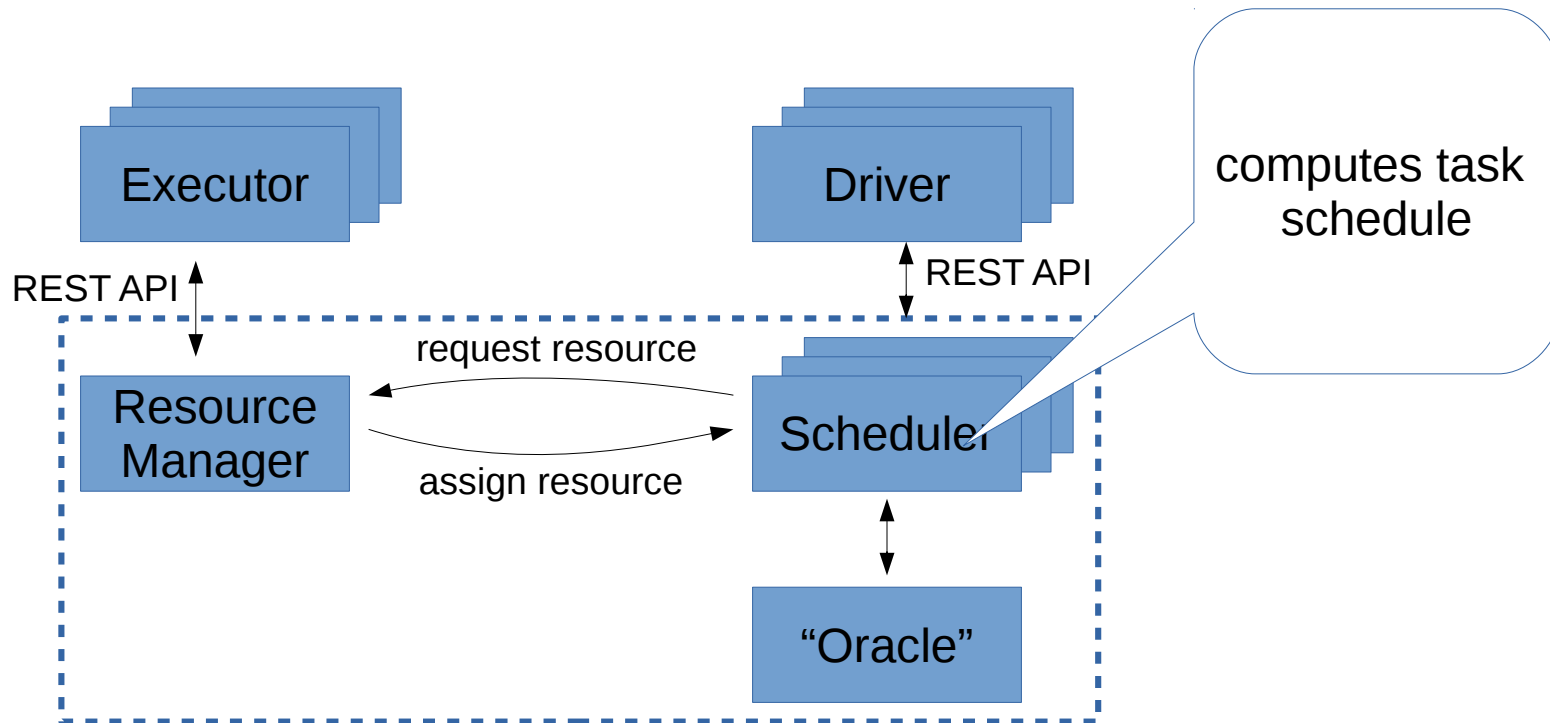
Architecture Overview




HCL Scheduler




HCL Scheduler



HCL Scheduler

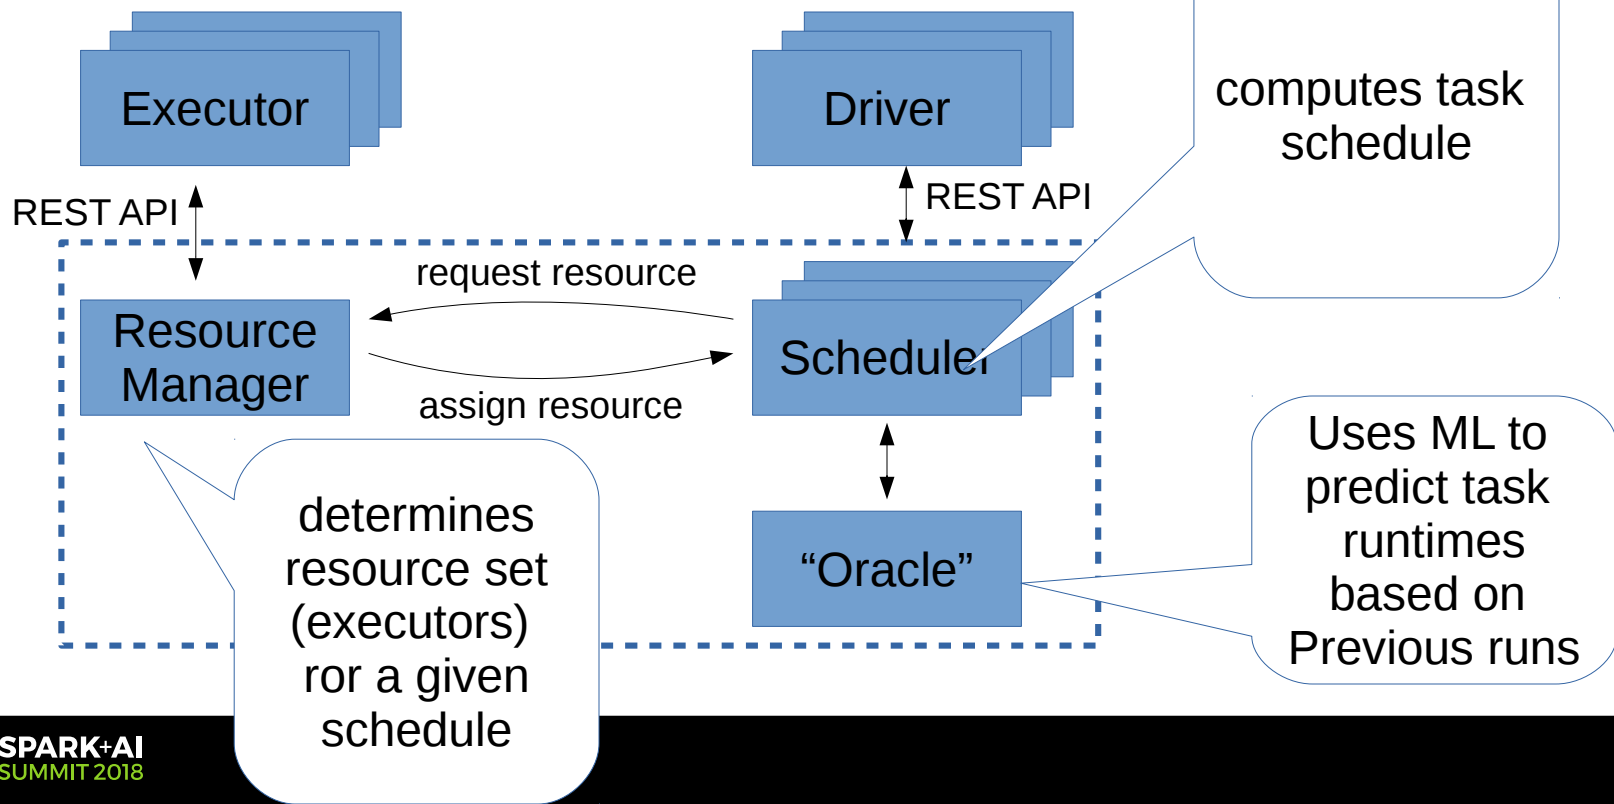


HCL Scheduler



HCL Scheduler

“The HCI Scheduler:
Going all-in on Heterogeneity”,
Michael Kaufmann et al., HotCloud’17



Video: Putting things together

Application 1:
GridSearch

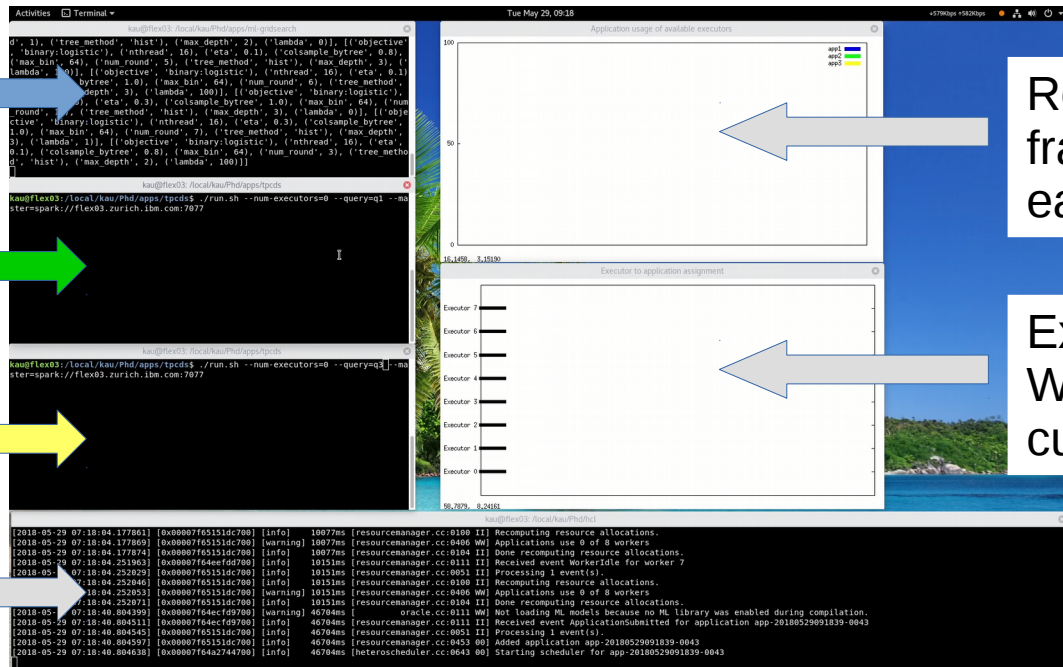
Application 2:
SQL TPC-DS

Application 3:
SQL TPC-DS

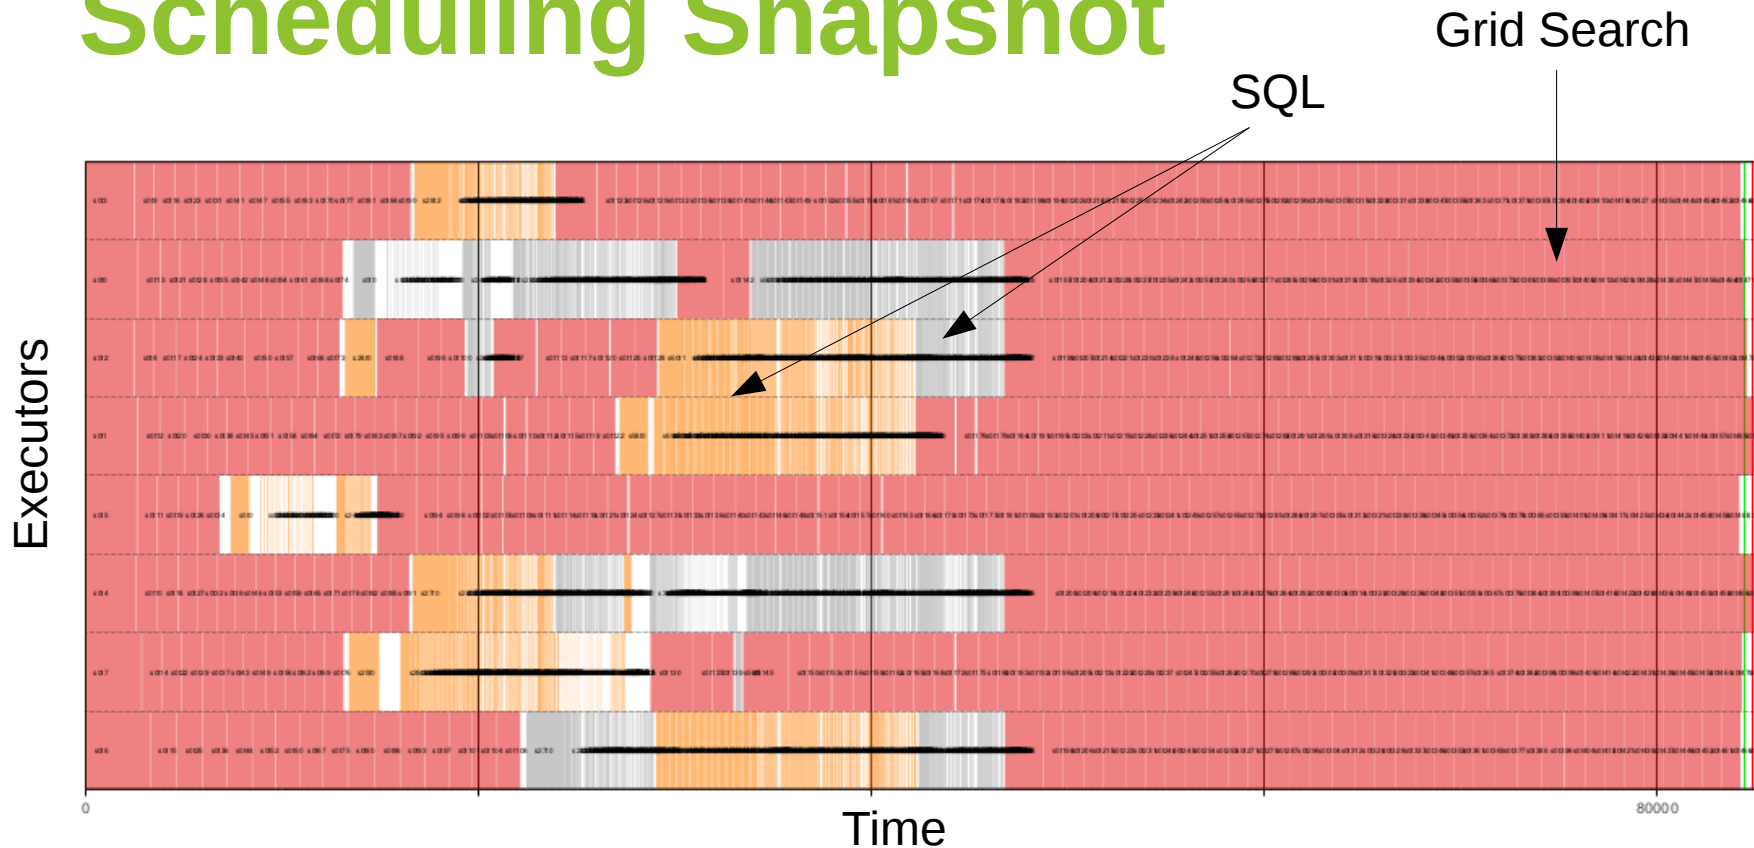
HCL
Scheduler

Resource view:
fraction of resources
each app consumes

Executor view:
Which app an executor
currently runs



Scheduling Snapshot



What about Performance?

- Workloads:
 - Deep learning (digit recognition) using Spark/Tensorflow, MNIST data set
 - SQL: TPC-DS
- Clusters:
 - 8 node cluster, 10Gb/s Ethernet
 - 8 node cluster, 100Gb/s RoCE
- Software
 - Spark2.3, Tensorflow 1.2

Conclusion

- Efficient serverless computing is challenging
 - Local state (e.g. shuffle, cached input, network state) is lost as compute cloud scales up/down
- This talk: turning Spark into a serverless framework by
 - Implementing a new serverless scheduler
 - Consequently storing compute state remotely using Apache Crail
- Supports arbitrary Spark workloads with almost no performance overhead
 - MapReduce, SQL, Iterative Machine Learning
- Implicit support for fast network and storage hardware
 - e.g, RDMA, NVMe

Future Work

- Add support for dynamic re-partitioning on scale events
- Add support for automatic caching
- Add more sophisticated scheduling policies

Thanks to

Michael Kaufmann, Adrian Schuepbach, Jonas Pfefferle,
Animesh Trivedi, Bernard Metzler, Ana Klimovic, Yawen
Wang

Backup

Workloads and Frameworks

	Microservices	Workflows	MapReduce	SQL	ML
AWS λ, Google CF, Azure F					
AWS λ + AWS StepFunction					
PyWren					
Databricks Serverless					

Serverless frameworks not designed to run arbitrary workloads