

# Serverless Machine Learning on Modern Hardware

IBM Research

**#Res6SAIS**

# Serverless Computing



- No need to setup/manage a cluster
- Automatic, dynamic and fine-grained scaling
- Sub-second billing
- AWS Lambda, Google Cloud Functions, Azure Functions, Databricks Serverless


# Challenge: Performance

- **Container startup:** may have to dynamically spin up containers per function call
  - Takes several 200-300 milliseconds for a “cold” container
- **Storage:** input data needs to be fetched from remote storage (e.g., S3 object store)
  - As opposed to compute-local storage, e.g., HDFS
- **Data sharing:** intermediate needs to be temporarily stored on remote storage (e.g. S3, Redis)
  - Becomes problematic as workloads get more complex
  - Affects operations like shuffle, broadcast, etc.,

# Challenge: Performance

- **Container startup:** may have to dynamically spin up containers per function call
  - Takes several 200-300 milliseconds for a “cold” container
- **Storage:** input data needs to be fetched from remote storage (e.g., S3 object store)
  - As opposed to compute-local storage, e.g., HDFS
- **Data sharing:** intermediate needs to be temporarily stored on remote storage (e.g. S3, Redis)
  - Becomes problematic as workloads get more complex
  - Affects operations like shuffle, broadcast, etc.,

# Example: MapReduce (Cluster)




data is mostly  
**written** and  
**read** locally

# Serverless MapReduce

Dynamically  
growing/shrinking  
compute cloud




Map  
Stage



Shuffle



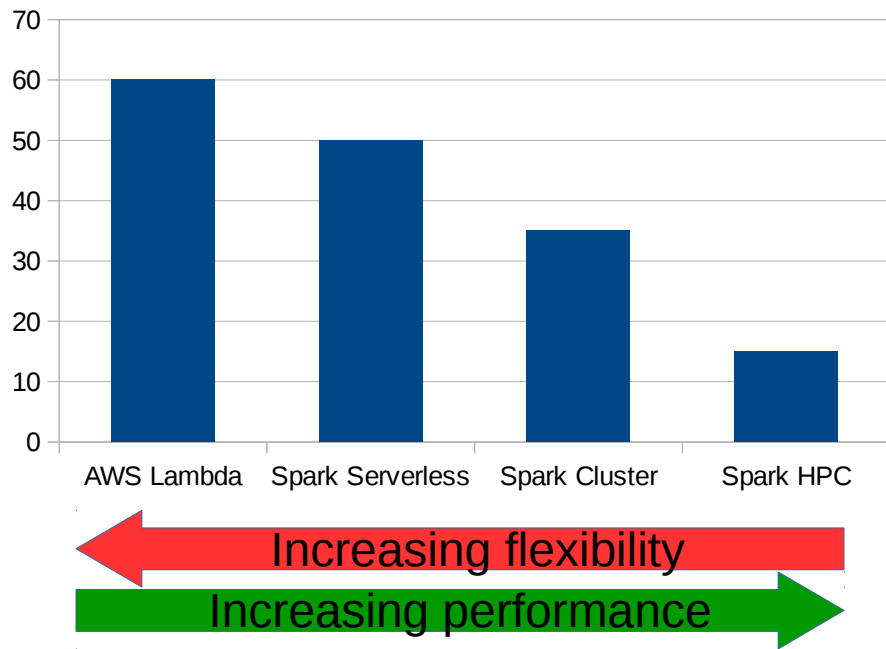
Reduce  
Stage



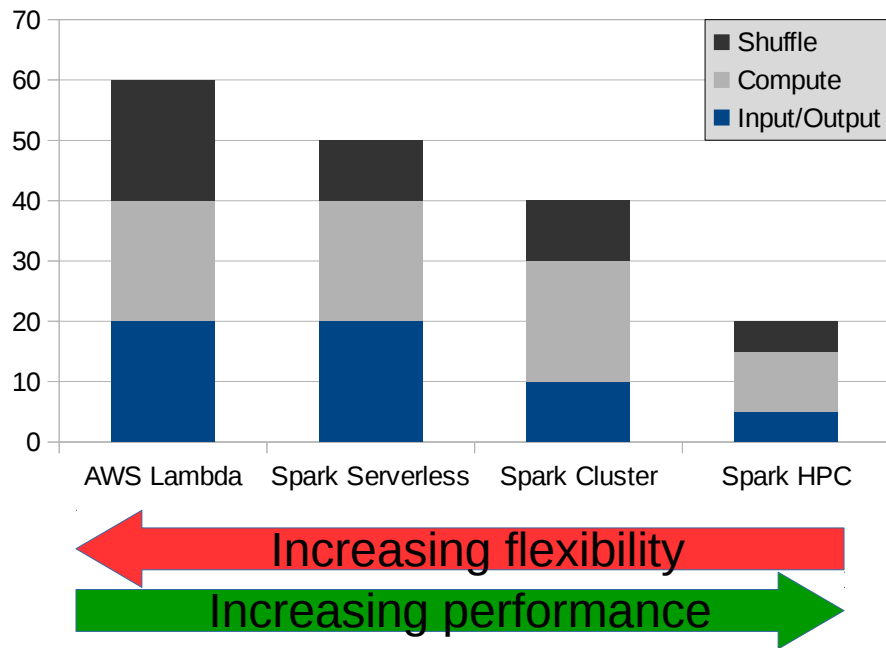
Storage Service  
(e.g, S3, Redis)

data is  
exclusively  
**written** and  
**read** remotely

# Sorting 100GB




# Is I/O a problem?





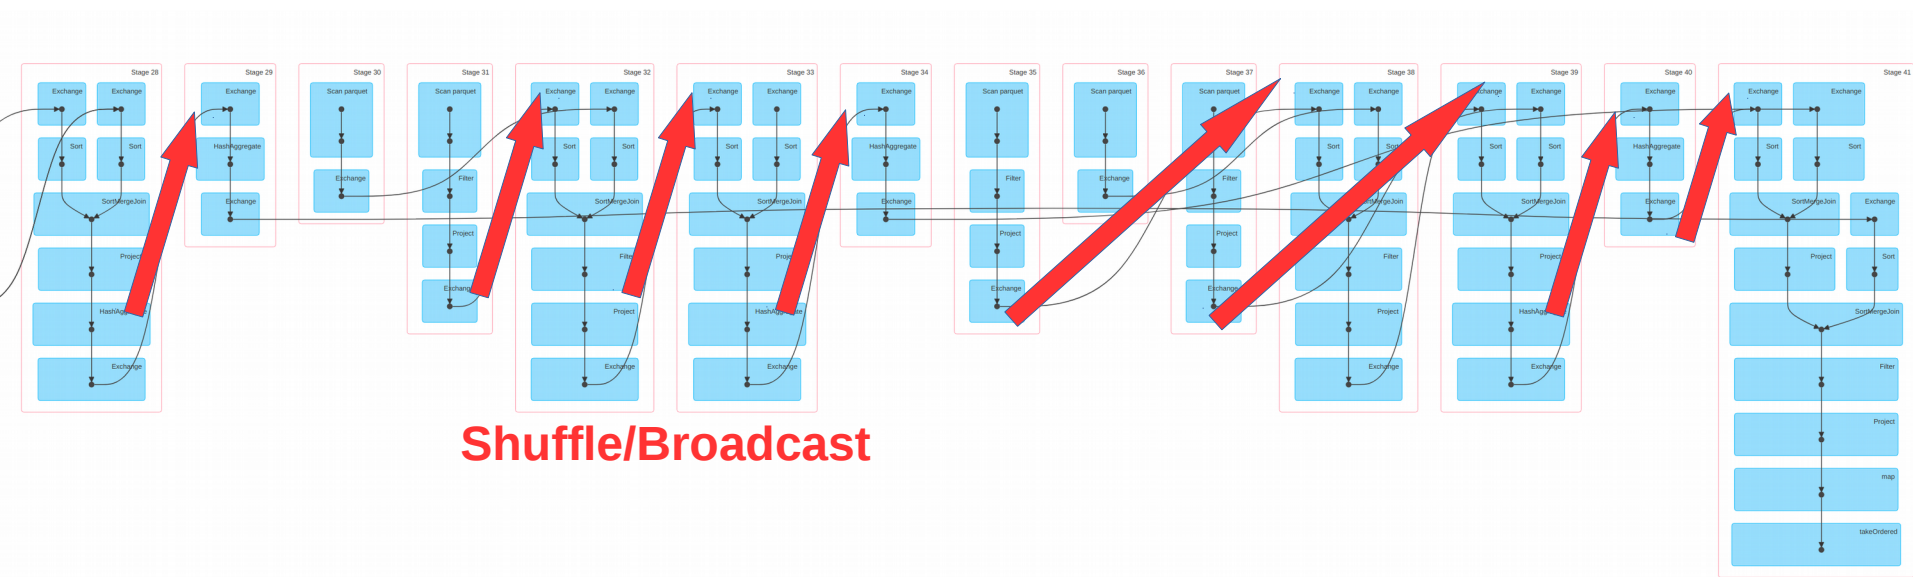
# What about other workloads?

Example: SQL, Query 77 / TPC-DS benchmark



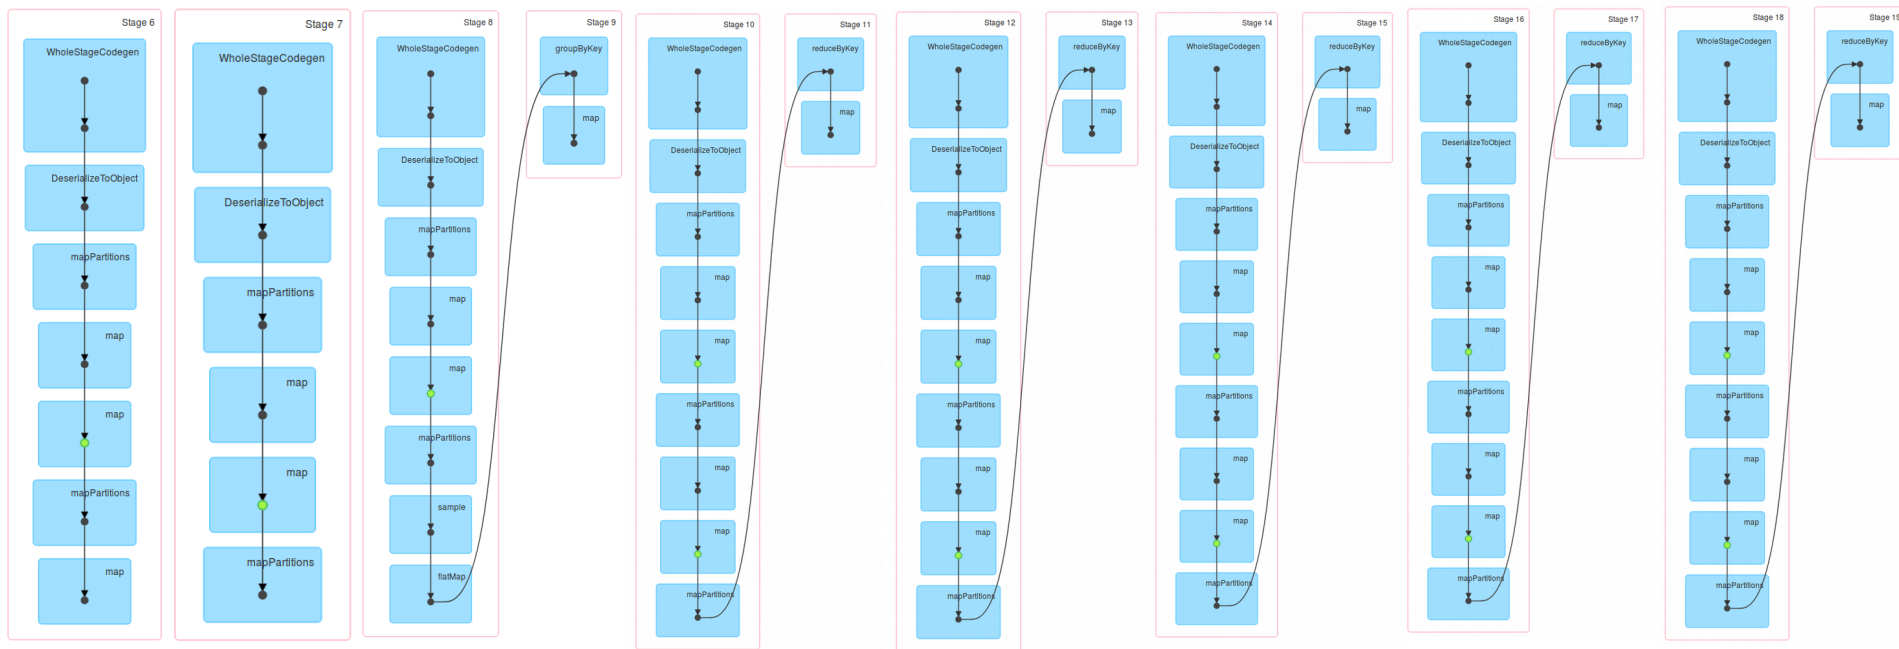
# What about other workloads?

Example: SQL, Query 77 / TPC-DS benchmark



# What about other workloads?

Example: RandomForest



# Workloads and Frameworks


	Microservices	Workflows	MapReduce	SQL	ML
AWS λ, Google CF, Azure F					
AWS λ + AWS StepFunction					
PyWren					
Databricks Serverless					

Serverless frameworks not designed to run arbitrary workloads

# Challenge #2: Workloads

- Serverless originally designed for simple use cases
  - E.g., image post-processing triggered by an upload
- What about more complex workloads?
  - MapReduce: Can be implemented on top of most frameworks
  - SQL? Databricks
  - Machine Learning?

# Serverless: Different Workloads



Serverless frameworks not designed to run arbitrary workloads

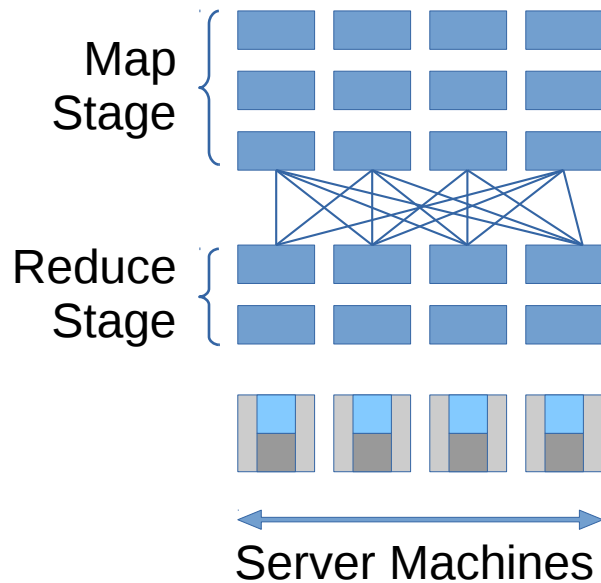
# Backup

# Template Tite

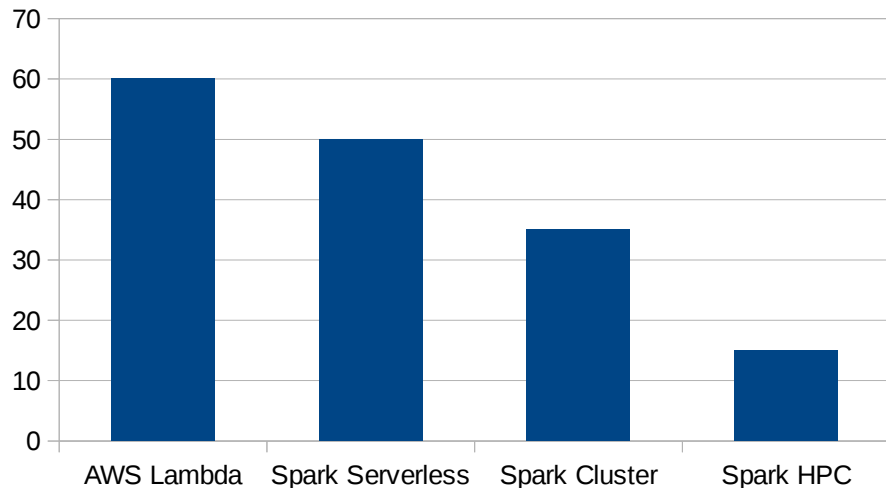
- Template List
- Template List
  - Template item



# Example: MapReduce




# Sorting 100GB



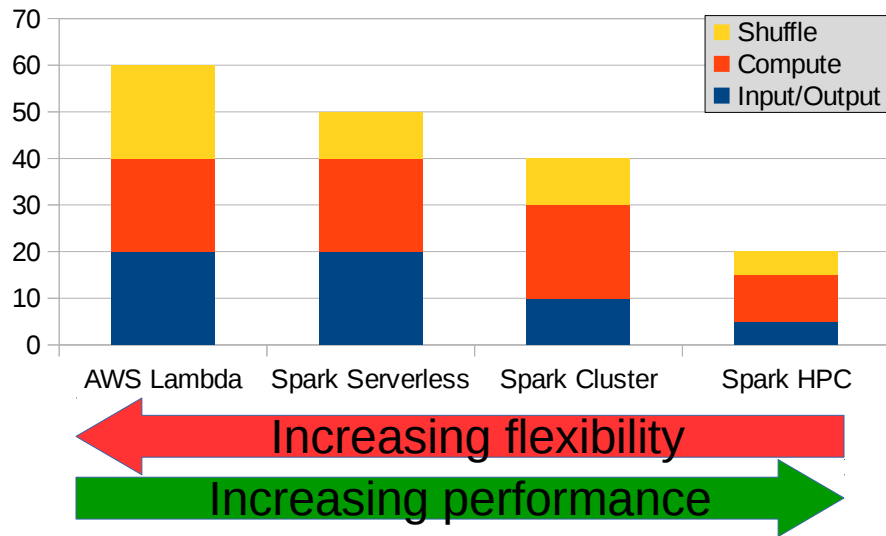
Serverless execution is 3-4x slower than an optimized cluster configuration

# Sorting 100GB



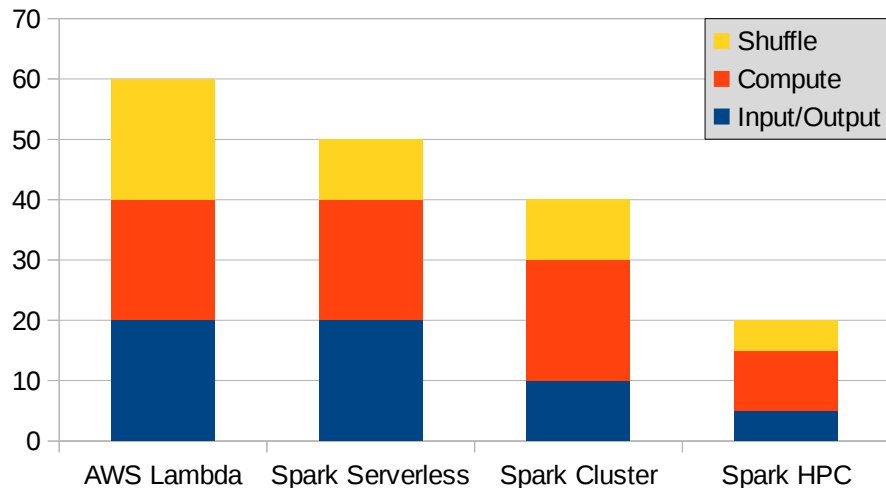
Serverless execution is 3-4x slower than an optimized cluster configuration

# Sorting: Is I/O a problem?



With serverless, substantial amount of time is spent on reading from remote storage

# Sorting: Is I/O a problem?



With serverless, substantial amount of time is spent on reading from remote storage

# What about other workloads?

Example: CoCoa...

